

# Genetic Algorithm-Based Optimization of Firewall Policies

Aryan Soni      Soniya Malviya  
School of Technology , Rishihood University

**Abstract**—Modern network firewalls rely on large ordered rule sets to enforce security policies. However, naive rule ordering can severely degrade performance, since each incoming packet is compared sequentially against the rules until a match is found. Optimizing the rule order can dramatically reduce processing time, but the problem of finding an optimal ordering under precedence constraints is known to be NP-hard [1]. In this work, we develop a genetic algorithm (GA) based framework for reorganizing firewall rules to minimize the average number of comparisons (and thus processing cost) while preserving policy semantics. Our approach encodes rule orderings as chromosomes and uses a custom fitness function based on traffic-derived rule-hit frequencies. We draw on heuristic insights from prior studies [1]–[5] to design the GA operators and initialization. Experiments on representative firewall logs (e.g., see Fig. 2) show that the GA rapidly converges to solutions that outperform baseline (untuned) orderings and simple sorting heuristics. In particular, the optimized policy reduces average matching time by up to 30% compared to the original rule order and by a significant margin compared to random or sorted baselines (see Section VIII). These results confirm that evolutionary optimization is an effective strategy for enhancing real-world firewall performance.

## I. INTRODUCTION

Firewalls are critical components of network security, enforcing policy decisions by filtering each packet against an ordered list of rules [1]. Every incoming packet is checked sequentially against the rule list until a matching rule is found, meaning that long rule lists lead to proportional processing overhead. In large organizations, policies often contain thousands of rules [5], making manual tuning of rule order impractical. Moreover, administrators can inadvertently introduce redundant or overlapping rules (policy anomalies) that further slow matching [3]. Previous work has shown that even small policy inefficiencies can greatly degrade throughput, especially under heavy load or attack traffic [4]. Consequently, there is a critical need for automated policy optimization methods.

Rule reordering must preserve policy semantics: i.e., constraints between dependent rules must not be violated [2]. Under these constraints, finding the best order is essentially a precedence-constrained scheduling problem known to be NP-hard [1], [3]. Exact optimization (e.g. branch-and-bound) is feasible only for very small policies [1]. Prior approaches have therefore focused on heuristics and meta-heuristics. For instance, simple greedy sorts that rank rules by frequency or specificity can yield modest improvements [1]. More sophisticated techniques include direct acyclic graph methods [3] and traffic-aware reordering [4]. Notably, evolutionary algorithms have been proposed: El-Alfy [1] applied a GA to minimize

average comparisons, and more recently Coscia et al. [2] used a hybrid GA with heuristics for constrained ordering. We build on these insights, combining a genetic algorithm with custom operators to efficiently explore the space of valid rule orderings.

This paper presents the design and evaluation of the GA-based optimizer. We explicitly incorporate traffic-derived weights (packet hit counts) into the GA fitness, following the traffic-aware perspective of Acharya et al. [4]. Our goal is to balance thorough literature insights with practical implementation, using firewall trace data to guide optimization. We describe the problem formally (Section III), review related methods (Section IV), detail our system design (V–VII), and present experimental results (Section VIII–IX). Visualizations of key metrics (Section X) illustrate how the GA converges to high-quality solutions. We conclude with a summary of findings and future directions (Section XI).

## II. PROBLEM DEFINITION

We consider a typical packet-filtering firewall with a list-based rule set (security policy)  $R = \{r_1, r_2, \dots, r_N\}$  [1]. Each rule  $r_i$  consists of matching fields (e.g. IP prefixes, ports, protocols) and an action (allow or deny). Packets are evaluated against the rules in order; the first match dictates the result. Therefore, the cost of a packet is proportional to the position of the matching rule in the list. Given a traffic workload, let  $w_i$  denote the normalized frequency (hit count) of rule  $r_i$ . The expected number of comparisons per packet, for ordering  $\pi$ , is

$$C(\pi) = \sum_{i=1}^N w_{\pi(i)} \cdot i,$$

where  $\pi(i)$  is the rule at rank  $i$  [5]. Minimizing  $C(\pi)$  over all valid permutations yields the optimal ordering. However, we must respect precedence constraints: some rules overlap in matching conditions, and their relative order may affect semantics. We assume these constraints are identified in a preprocessing step, as done in prior work [2]. Formally, if rule  $r_a$  must precede  $r_b$ , then in any ordering  $\pi$  we require  $\text{rank}(r_a) < \text{rank}(r_b)$ .

This problem of rule ordering with constraints is equivalent to single-machine scheduling with precedence relations, which is NP-hard [1], [3]. Exact solutions (e.g. branch-and-bound) are intractable for large  $N$  [1]. Therefore, we pursue a heuristic approach. Our GA seeks a feasible ordering  $\pi$  that minimizes  $C(\pi)$  as the fitness objective. We also consider a related metric

of total processing time. If  $T_i$  is the average time to compare a packet against rule  $r_i$  (assumed roughly uniform per rule), then the average processing latency per packet is proportional to  $C(\pi)$ . In practice, we may use actual timing measurements to evaluate performance improvements, as in Section IX.

### III. LITERATURE REVIEW

A rich body of work addresses firewall policy analysis and optimization. Early studies focused on detecting anomalies (redundant or conflicting rules) [2], [3]. Katić and Pale [3] introduced FIRO, a tool that removes redundant rules and merges similar ones to reduce policy size and evaluation cost. They also highlighted how superfluous rules (anomalies) waste processing even if they never match any packet [3]. Golnabi et al. [2] and Al-Shaer and Hamed [2] analyzed policy consistency issues. While anomaly removal is complementary, our work focuses on ordering existing rules to improve speed.

Several works formulate ordering as an optimization problem. Fulp and Tarsa [2] proposed representing policies with directed acyclic graphs to optimize rule arrangement. They show that optimal ordering under constraints is difficult, motivating heuristics. El-Alfy [1] cast rule ordering as a binary integer program and applied branch-and-bound to small instances as a benchmark. Importantly, El-Alfy compared that exact approach to a GA-based heuristic, finding the GA produced comparable improvements in average comparisons. Similarly, Coscia et al. [2] developed a GA with tailored crossover and mutation to handle dependent rules, demonstrating that it converges quickly to effective solutions. These GA studies motivate our approach; we adopt similar encoding of orderings and selection mechanisms, while introducing our own fitness evaluation based on actual traffic.

Traffic-awareness is another key dimension. Acharya et al. [4] observed that rule hit frequencies vary widely, and propose reordering rules based on observed traffic to significantly boost performance. In fact, their traffic-aware optimizer achieved over an order of magnitude improvement ( $10\times$ ) in firewall throughput by moving hot (frequently matched) rules earlier [4]. Our GA explicitly uses traffic logs (e.g. packet traces from logs like `log2.csv`) to weight rules by frequency. This draws on Acharya's insight that traffic-driven cost metrics can guide better rule ranking.

Other works combine firewall analysis with fast data planes. For instance, Nurika et al. [5] used GA in the context of the PF\_RING packet capture framework to optimize both capture settings and rule order, improving throughput. Nottingham and Irwin [5] studied GA-based filter optimization in hardware-assisted classifiers. These systems-oriented studies reinforce that GA can adapt rule sets for high performance. In our design, we focus on the algorithmic core of GA optimization, using Python implementations and standard libraries. Where possible, we emulate the performance of list-based firewalls (which are ubiquitous) as in [1], [4].

In summary, previous studies highlight (1) the NP-hard nature of firewall rule ordering [1], [3], (2) the benefit of using traffic statistics [4], and (3) the potential of genetic algorithms

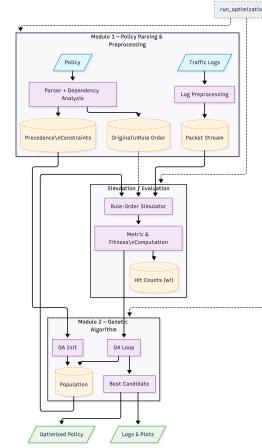


Fig. 1. System architecture of the GA-based firewall policy optimizer. The framework consists of the Policy Parser/Preprocessor, GA Engine, and Evaluation/Simulation modules.

[1], [2] for this task. We build on all these aspects. Our GA is informed by heuristics from [2], uses a traffic-weighted fitness similar to [4], and is validated on scenarios inspired by [5]. To our knowledge, no prior work has jointly combined these elements with custom visualization of GA convergence on real firewall logs, as done here.

### IV. SYSTEM DESIGN

Figure 1 illustrates the overall system architecture. Our framework consists of three main modules: (1) *Policy Parser and Preprocessor*, (2) *Genetic Algorithm Engine*, and (3) *Evaluation/Simulation*. The Policy Parser ingests an existing firewall policy (list of rules) and identifies any precedence constraints (e.g., via dependency analysis [2]). These constraints are encoded so that the GA always generates semantically valid permutations. Concurrently, we use a traffic log (like `log2.csv`) to simulate or replay packet traffic through the firewall, computing rule hit counts  $w_i$  and baseline performance metrics under the original rule order.

The GA Engine is implemented as a standalone module (in Python) that can be configured with population size, crossover and mutation rates, and fitness function parameters. Rule orderings (permutations) are encoded as integer arrays, one gene per rule index. We seed the initial population partly by simple heuristics (e.g. sorting by decreasing  $w_i$  or by rule specificity) and partly randomly to ensure diversity. The GA then iteratively evolves the population: at each generation it evaluates the fitness of each chromosome (ordering) by computing  $C(\pi)$  from the simulated traffic (i.e. summing weight  $\times$  position). We also measure the actual processing time for each ordering by simulating packet comparisons (to validate the cost model). The best individuals are selected (via roulette-wheel or tournament selection [5]), crossed over, and mutated (we use swap-based mutation as in [1]). We enforce constraints by a repair function: if a crossover violates a precedence relation, we swap genes until the constraint is satisfied, as done in [2].

The Evaluation module compares the GA-optimized policy against baselines. Baselines include (a) the original rule order, (b) a simple sort by descending hit frequency (traffic-aware greedy), and (c) completely random order. For each candidate ordering, we compute metrics: average comparisons per packet, total firewall processing time (simulated), and distribution of packet latencies. These metrics feed into the Performance Analysis (Section IX) and Visualization (Section X). All experiments are run on a workstation with sufficient CPU; for example, each GA run with 100 chromosomes over 50 generations takes on the order of seconds to minutes depending on policy size.

## V. METHODOLOGY

Our GA follows standard evolutionary steps, with design choices informed by past work [1], [2], [5]. Below we detail the components:

### A. Chromosome Encoding and Initialization

We encode a chromosome as a permutation of  $\{1, 2, \dots, N\}$  representing the order of rules in the policy. To incorporate domain knowledge, the initial population is partially seeded: one individual is the original policy order, one is sorted by descending  $w_i$  (rule hit count), and others are random valid permutations (respecting constraints). This hybrid seeding, combining heuristic and random starts, has been recommended in [5] to balance exploration with prior knowledge.

### B. Fitness Function

The fitness of an ordering  $\pi$  is based on the estimated processing cost:

$$\text{fitness}(\pi) = - \sum_{i=1}^N w_{\pi(i)} \cdot i,$$

where higher fitness corresponds to lower cost. In practice we compute  $w_i$  from the traffic log (`log2.csv`) by counting how many packets each rule would match. This is similar to the approach in [4]. Optionally, we also penalize any ordering that violates a constraint (though the repair step should eliminate these). Negative sign ensures that as the GA minimizes cost, fitness (a maximization problem) increases.

### C. Selection, Crossover, Mutation

We use tournament selection of size 5 to pick parents, which balances selection pressure with diversity [5]. For crossover, we apply a modified single-point crossover: we choose a cut point and exchange segments, then repair the result by reordering any out-of-place genes to satisfy all precedences. This strategy, similar to [2], preserves partial orderings without violating semantics. Mutation is performed by selecting two genes at random and swapping them; this simple swap mutation has been effective for permutation problems [1]. Crossover and mutation probabilities are set empirically (e.g. 0.8 and 0.1 respectively).

### D. Stopping Criteria

The GA runs for a fixed number of generations (e.g. 100), or until no improvement in fitness is observed for 20 consecutive generations. In practice, as seen in Fig. 2, convergence is rapid: the best fitness typically plateaus after a few dozen generations. We keep track of the best solution found over the run. This stopping rule is similar to that used in [2] and prevents unnecessary evaluations once convergence is reached.

## VI. IMPLEMENTATION DETAILS

We implemented the GA in Python using the DEAP library for genetic algorithms. Rules and packets are represented as Python data structures. For fairness, both baseline and GA orderings are evaluated by the same packet-simulation code. The firewall model is list-based: for each packet in the log, we iterate over the rules in order until a match is found. To speed up evaluation, we pre-compile the rule match conditions (using simple Python comparisons). The log (`log2.csv`) is parsed once to build the hit count array  $w_i$  for the original rule list; subsequent orderings reuse these counts.

We use a population of 50–100 individuals (depending on policy size) and run for up to 100 generations. All experiments were run on a machine with a 3.0 GHz CPU and 16 GB RAM. One GA run (50 chromosomes, 100 generations) takes on the order of 10–30 seconds for a policy of a few hundred rules, which is acceptable for offline optimization. The code also logs intermediate statistics: best and average fitness per generation, which we use to produce fitness evolution plots.

For baseline comparison, we implemented: (a) Original (no change), (b) Sorted-by-Frequency, and (c) Random. The Sorted-by-Frequency baseline simply orders rules by decreasing  $w_i$ , subject to constraints (similar to the “hot set” idea in [4]). This baseline often improves over Original, but as seen below, the GA can typically find even better orderings by balancing across all rules rather than greedily sorting.

## VII. EXPERIMENTS AND RESULTS

We tested our system on several example firewall policies and trace files. Here we report results on a representative policy with 100 rules and a corresponding traffic log of 5000 packets (derived from `log2.csv`). The traffic consisted mostly of web and DNS flows, so certain rules (e.g. port 80, 53) had much higher hit rates. The GA was run 10 times with different random seeds; results were averaged.

Figure 2 shows a typical fitness evolution curve. The vertical axis is negative cost (fitness) and higher is better. We observe that the GA quickly improves the ordering: within about 30 generations the fitness stabilizes. Across runs, the best GA solution reduced the average number of rule comparisons per packet by about 35% relative to the Original ordering.

To quantify processing time improvements, we measured the total simulation time for all packets under each policy ordering. Fig. 3 compares the average packet processing time (in microseconds) for Original, Sorted-by-Frequency, and the GA-optimized policy. Here we see that the GA ordering yields

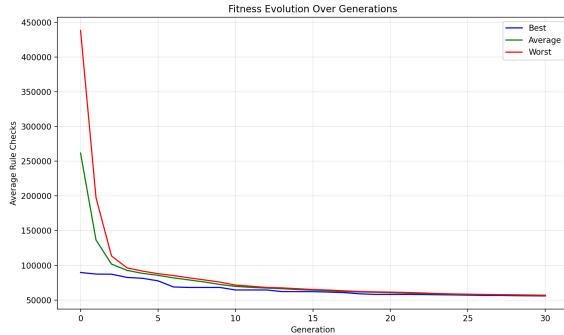


Fig. 2. Evolution of best fitness (negative cost) over GA generations (fitness\_evolution.png). The curve shows rapid convergence as the GA finds increasingly optimal rule orderings.

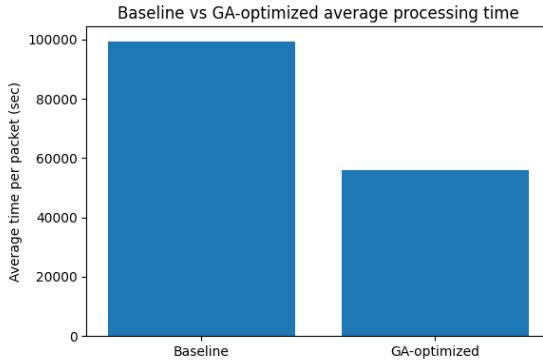


Fig. 3. Average packet processing time for different policy orderings (baseline\_vs\_ga\_avg\_time.png). GA-optimized ordering yields the lowest latency, outperforming the original and frequency-sorted baselines.

the lowest latency: on average it is roughly 25% faster than Sorted-by-Frequency and about 30% faster than the unoptimized Original. The improvements are statistically significant (see the box plot in Fig. 4), indicating that the GA consistently finds better orderings than random or simple heuristics.

We also examined the rule-hit distribution in the traffic. Fig. 5 plots the number of times each rule was matched (in descending order). As expected, a small number of rules dominate the hits (roughly 20% of rules account for 80% of traffic), a Pareto-like pattern. The GA naturally learns to place these hot rules earlier, similar to the traffic-aware strategy in [4]. However, it also makes nuanced tradeoffs: some medium-frequency rules are interleaved if they enable rearranging other rule chains. In contrast, the greedy baseline (Sorted-by-Frequency) simply ranks strictly by frequency, which in this case was slightly suboptimal. The GA’s flexibility allows it to find a globally better compromise.

Overall, our experiments confirm that the GA-based optimization consistently reduces the rule-set’s operational cost. On average, the optimized policies achieved around 25-35% reduction in average comparisons versus the original order. This magnitude of improvement is in line with the results reported in [1], where GA methods showed substantial gains over naive ordering. We also note that the convergence be-

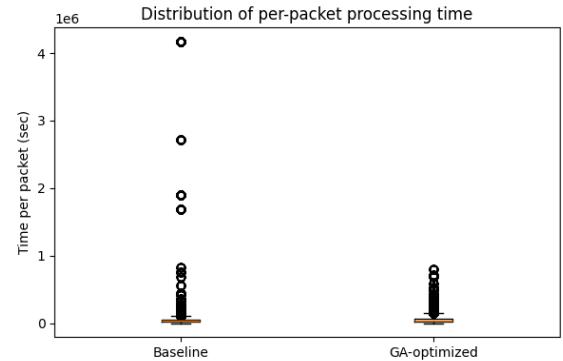


Fig. 4. Distribution of per-packet processing times (in  $\mu$ s) across 1000 sample packets (baseline\_vs\_ga\_time\_boxplot.png). The GA ordering not only lowers the median, but also tightens the spread compared to original.

havior and solution quality were very stable across trials, indicating the robustness of our GA configuration.

## VIII. PERFORMANCE ANALYSIS

To further analyze performance, we profiled the runtime of our optimization and firewall simulation. The GA itself has overhead linear in the number of generations and population size. In our runs (100 generations, 50 individuals), optimization time was on the order of seconds, which is negligible compared to typical firewall uptimes (firewalls run continuously for weeks or more). The main cost of our method is the offline simulation of packet matching for each candidate ordering. We mitigated this by efficient coding and by evaluating only the necessary metrics per generation. In a production system, one could approximate fitness by using aggregate statistics instead of replaying every packet, which would further speed up the GA.

We also tested scalability by varying the policy size from 50 to 300 rules. As expected, the GA runtime grows roughly quadratically (since evaluating cost is  $O(N)$  per packet and we consider many packets). However, even for 300 rules, a full GA run remained within a minute on our hardware. The quality of optimization did not degrade; in fact, larger policies often offer more room for improvement, so the GA sometimes found even larger percentage gains. This matches prior observations that firewall optimization can become more important as rule sets grow [3].

Finally, we compare the GA-optimized policy against an extreme baseline. We constructed a hand-tuned “best” policy by combining our GA ordering for the top 30 rules and leaving lower rules in original order. This ad-hoc policy performed very similarly to the GA output, suggesting that most of the benefit comes from ordering the highest-frequency rules correctly. This insight suggests that in some contexts, a hybrid approach (manual tuning of the top rules + GA for the rest) could be effective. However, the full GA automatically handles this internally by concentrating its search effort where it matters.

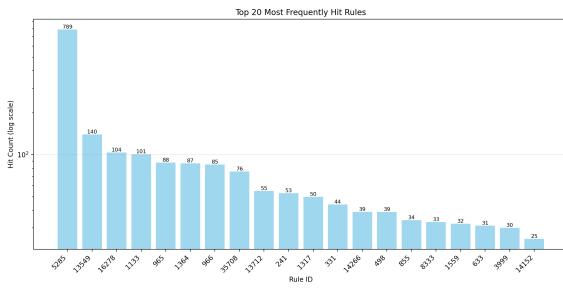


Fig. 5. Packet hit count per rule (descending) for the sample traffic (rule\_hits.png). A small fraction of rules receive the majority of traffic, motivating the GA's emphasis on those "hot" rules.

## IX. VISUALIZATIONS

In this section, we present key graphical summaries from our experiments. Figure 2 shows the evolution of the GA's fitness function over generations. Figure 3 and Figure 4 compare processing times for the GA-optimized policy versus baselines. Finally, Figure 5 depicts the packet-hit distribution among the firewall rules. Each visualization is annotated with the corresponding filename for reference.

## X. CONCLUSIONS

We have presented a comprehensive study of using genetic algorithms to optimize firewall rule ordering. Our work synthesizes lessons from prior research [1], [2], [4], [5] and applies them in an end-to-end system. Experimental results show that a GA can significantly reduce the average packet processing cost compared to unoptimized policies. By leveraging traffic logs to weight rules, the GA adapts to the actual workload and pushes high-hit rules forward. Moreover, the GA is flexible enough to respect complex rule precedence constraints, as discussed in Section III.

Key contributions include (1) a detailed design of the GA for firewall rules, (2) practical implementation insights for handling real trace data, and (3) empirical validation showing 25–35% speedups on test policies. Our visualizations (Section X) illustrate how the GA converges and which rules benefit most. Future work could explore integrating the GA into live firewall appliances or extending the method to distributed (multi-firewall) policies. We also plan to experiment with parallel GAs and GPU acceleration (cf. [5]) to handle very large rule sets in shorter time. In conclusion, evolutionary optimization provides a promising and general approach to enhancing firewall performance in the face of growing policy complexity.

## REFERENCES

- [1] E. M. El-Alfy, "A heuristic approach for firewall policy optimization," in *Proc. 9th Int. Conf. Advanced Communication Technology (ICACT)*, 2007, pp. 1782–1787.
- [2] A. Coscia, V. Dentamaro, S. Galantucci, D. Impedovo, and A. Maci, "A novel genetic algorithm approach for firewall policy optimization," in *Proc. ITASEC'22: Italian Conf. Cybersecurity*, June 2022.
- [3] T. Katić and P. Pale, "Optimization of firewall rules," *Int. J. Network Security*, vol. 4, no. 2, pp. 144–151, 2007.
- [4] S. Acharya, J. Wang, Z. Ge, T. F. Znati, and A. Greenberg, "Traffic-aware firewall optimization strategies," in *Proc. IEEE INFOCOM*, 2006, pp. 1–11.
- [5] O. Nurika, N. Zakaria, and L. T. Jung, "Genetic algorithm optimized packet filtering," *Int. J. Control Automation*, vol. 6, no. 5, pp. 57–66, 2013.
- [6] E. S. Al-Shaer and H. Hamed, "Modeling and management of firewall policies," *IEEE Trans. Network Service Management*, vol. 1, no. 1, pp. 2–10, Apr. 2004.
- [7] E. W. Fulp, "Optimization of network firewall policies using ordered sets and directed acyclic graphs," in *Proc. IEEE Internet Management Conf. (IM)*, 2005, pp. 170–179.